

Figure 1 — FAS vs flat retrieval: discrimination performance

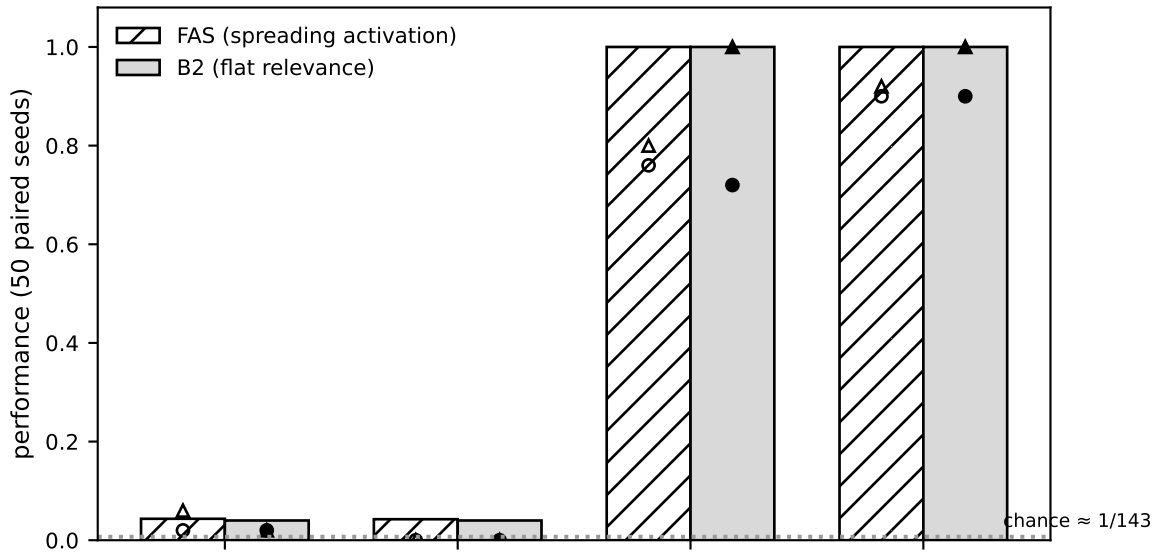


Figure 1. Confirmatory subset C-1...C-4 (50 paired seeds, C-2 LLM). Bars: median MRR; $\circ$ mean Recall@1 (outline=FAS); $\triangle$ mean Recall@3 (filled=FAS). Floor at $k=3$ inputs and ceiling at $k=5$ for both mechanisms; FAS is MRR-identical at ceiling (ns) and marginally higher at floor (S_C-1 raw $p=.006$, med diff = $+.0035$; below chance sensitivity for practical discrimination). B3 random baseline omitted ($\approx$ chance).